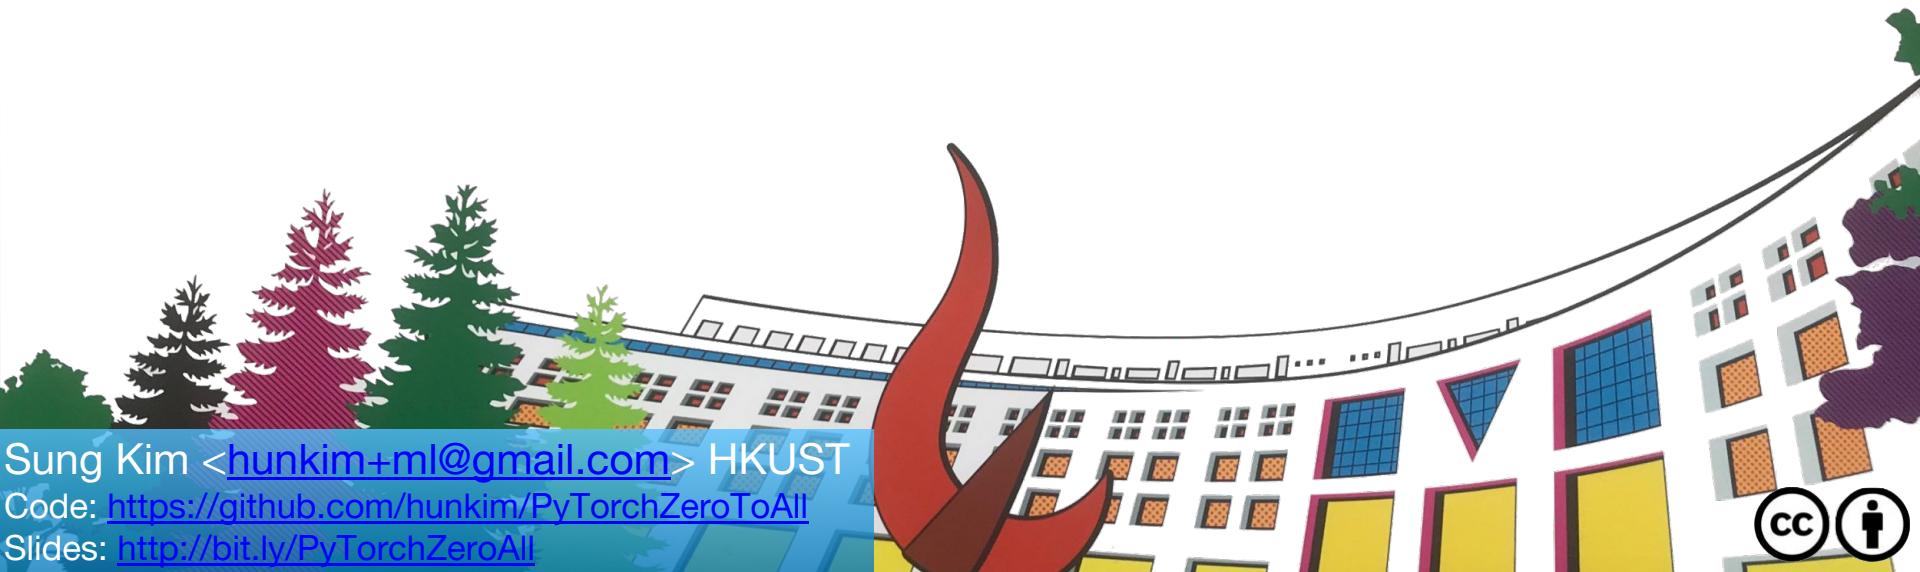


# ML/DL for Everyone with PYTORCH

## Lecture 10: CNN



Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST  
Code: <https://github.com/hunkim/PyTorchZeroToAll>  
Slides: <http://bit.ly/PyTorchZeroAll>



# Call for Comments

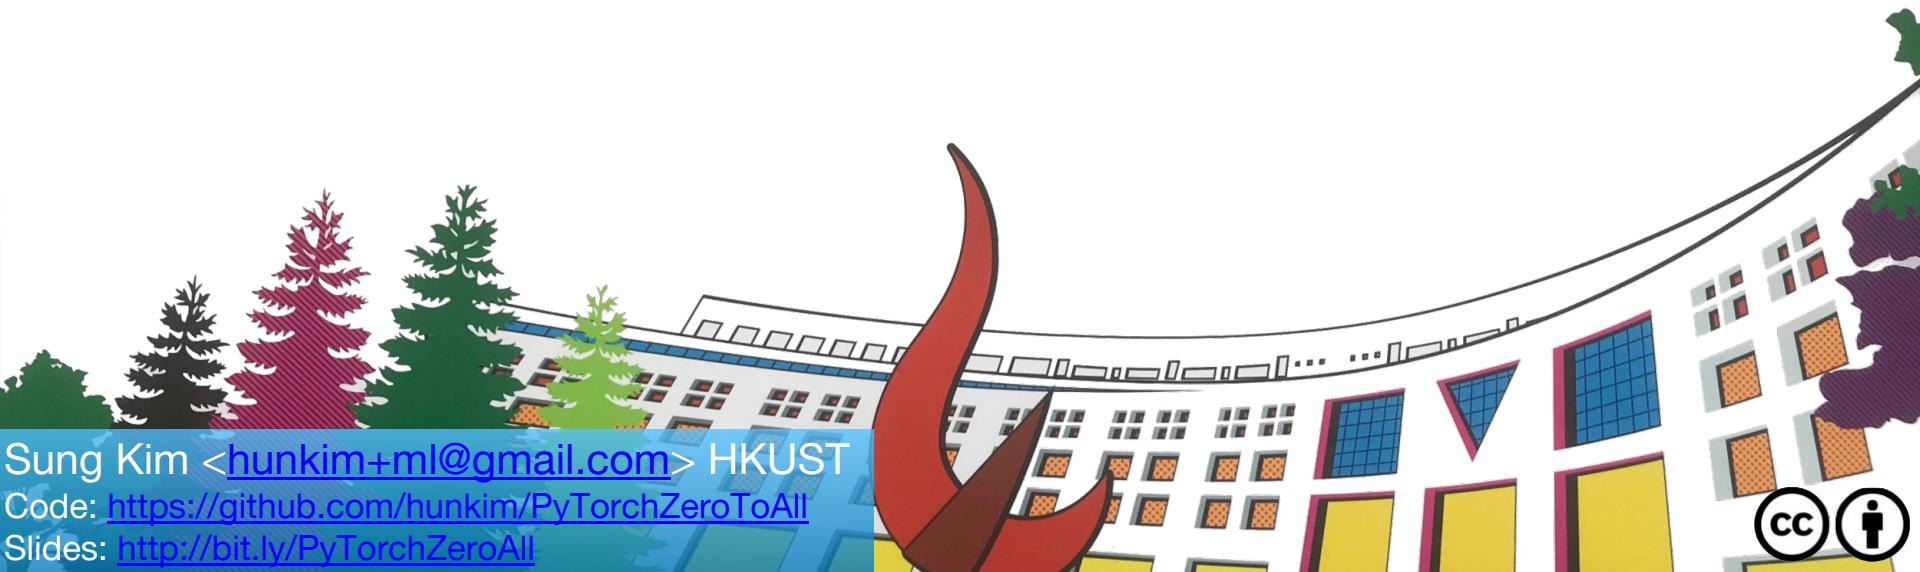
Please feel free to add comments directly on these slides.

Other slides: <http://bit.ly/PyTorchZeroAll>



# ML/DL for Everyone with PYTORCH

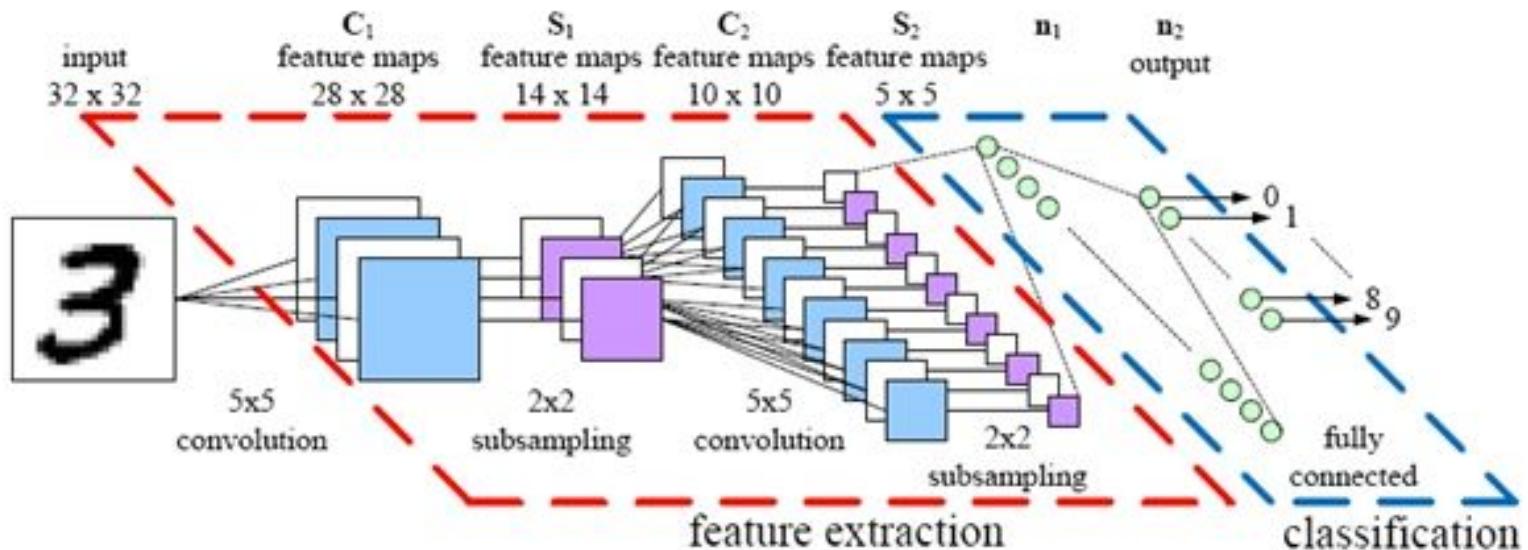
## Lecture 10: CNN



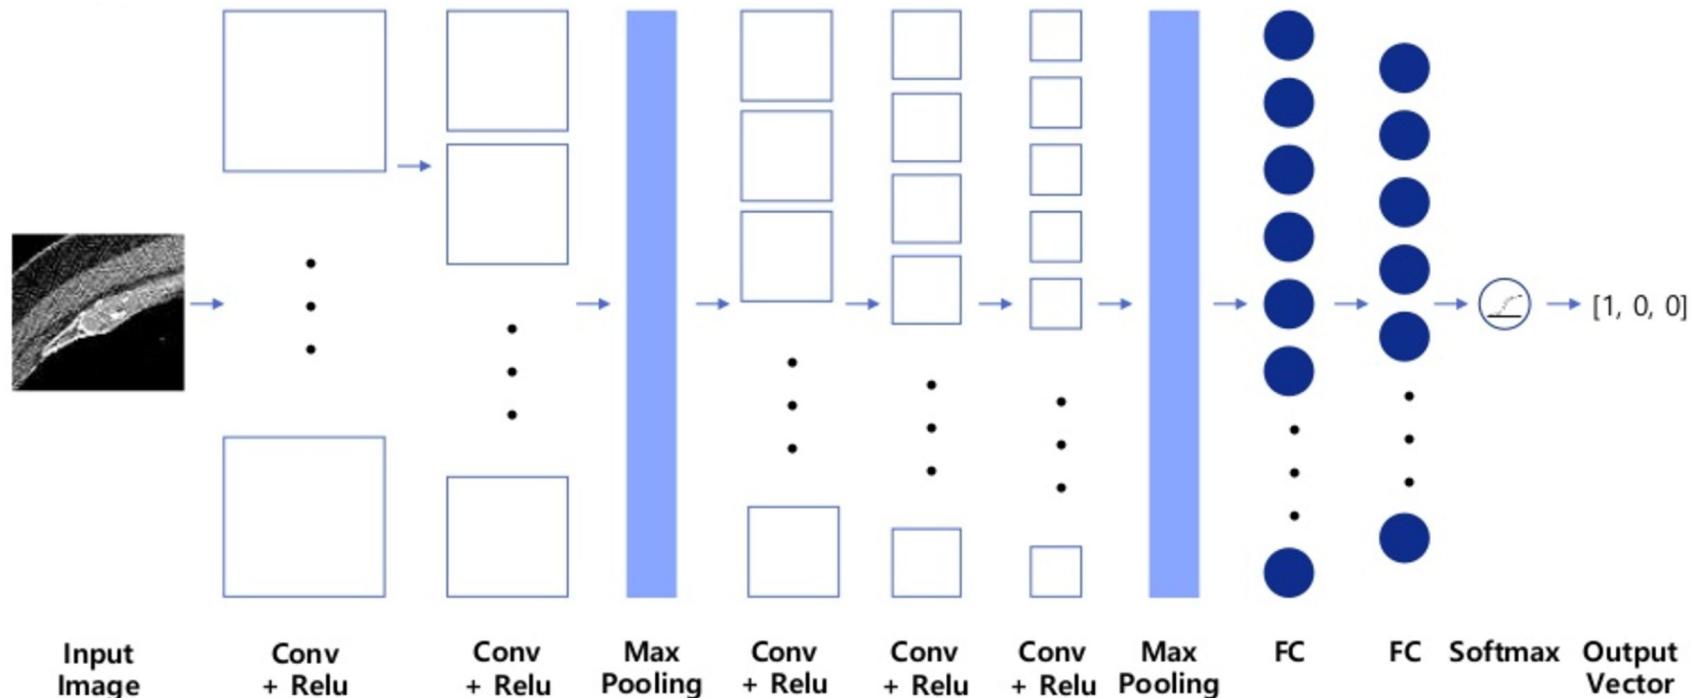
Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)> HKUST  
Code: <https://github.com/hunkim/PyTorchZeroToAll>  
Slides: <http://bit.ly/PyTorchZeroAll>



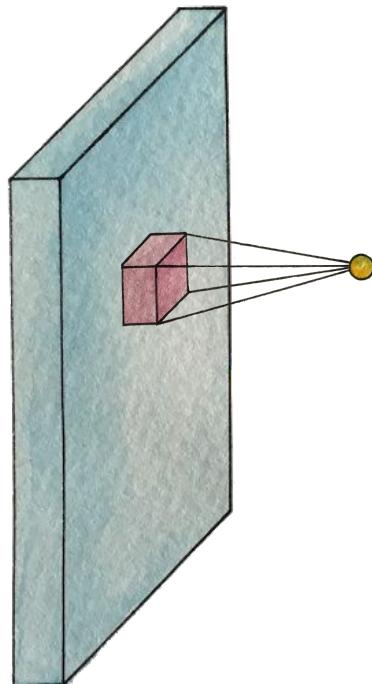
# CNN



# CNN for CT images



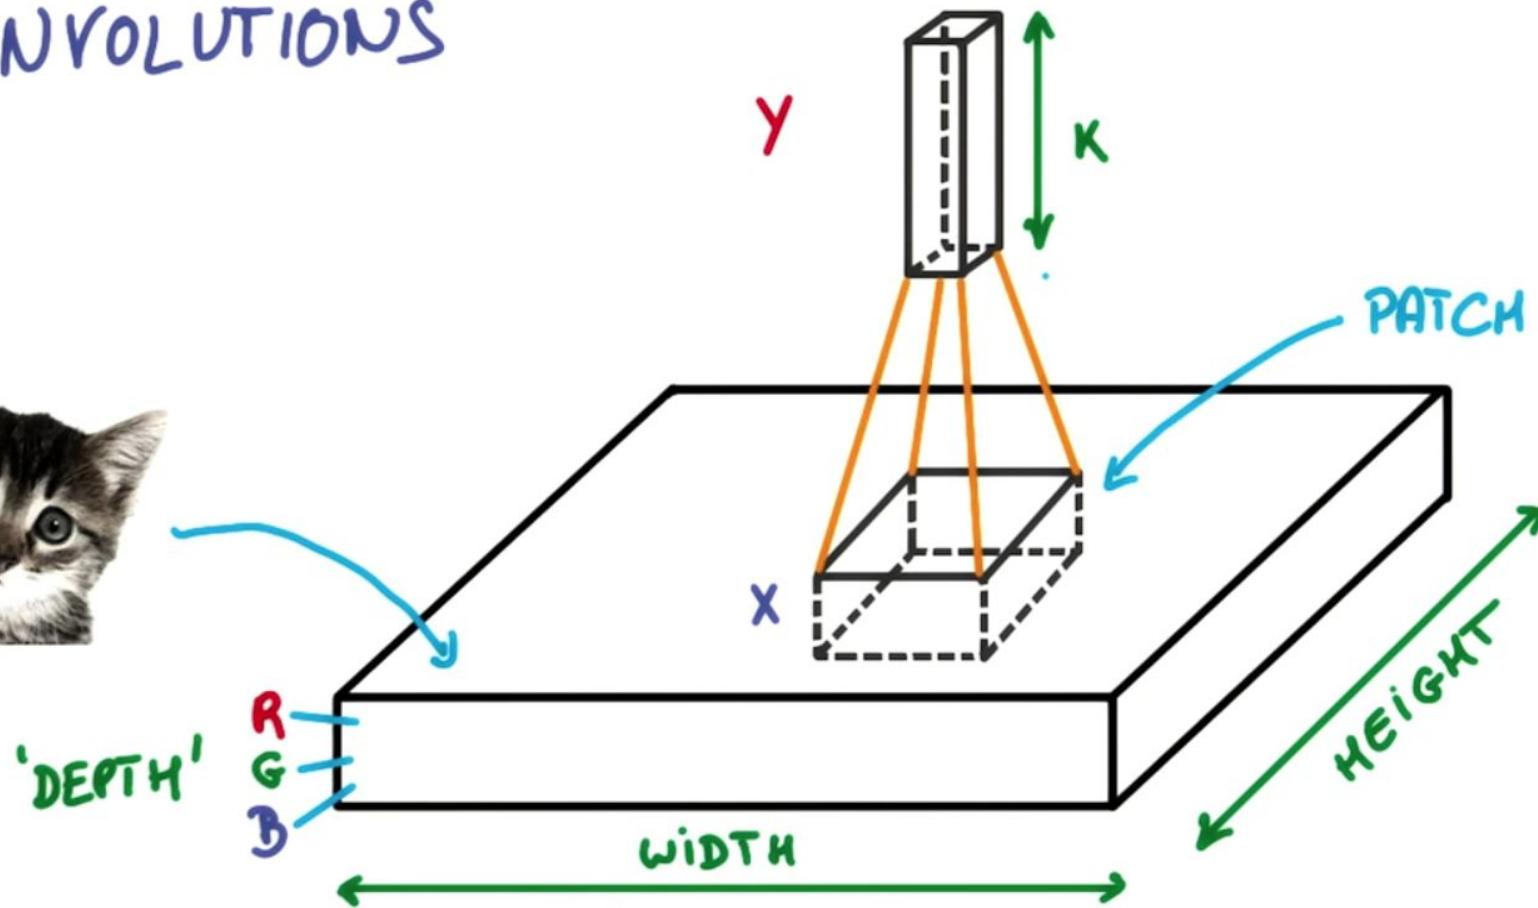
# Convolution layer and max pooling



Single depth slice

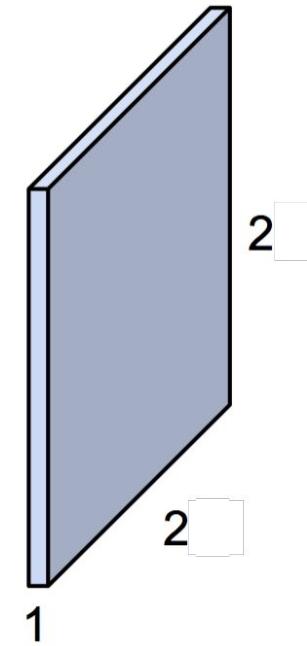
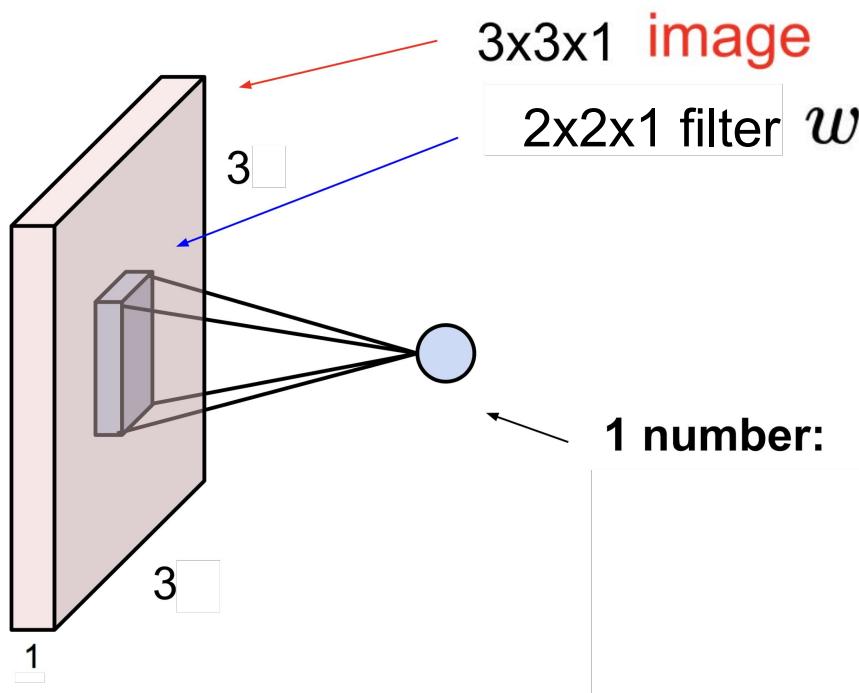
1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

# CONVOLUTIONS



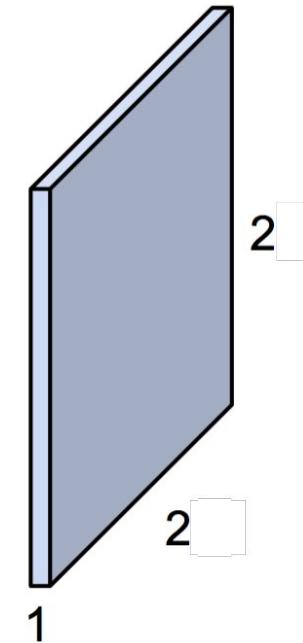
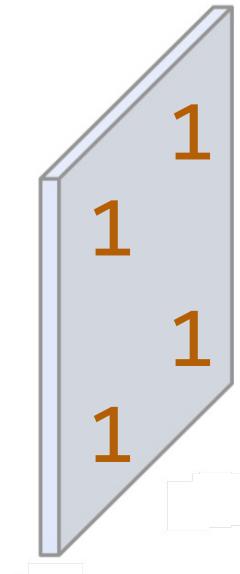
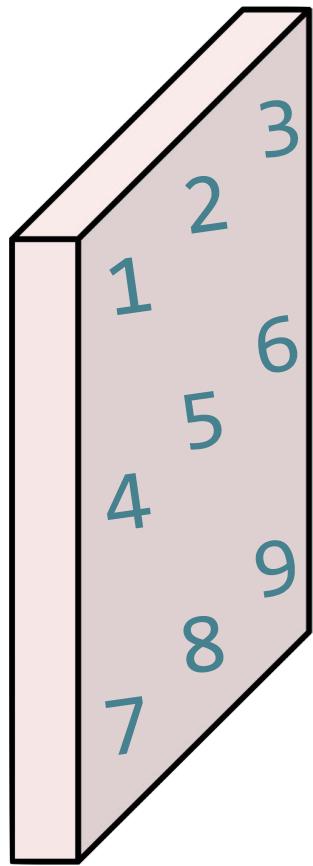
# Simple convolution layer

Stride: 1x1



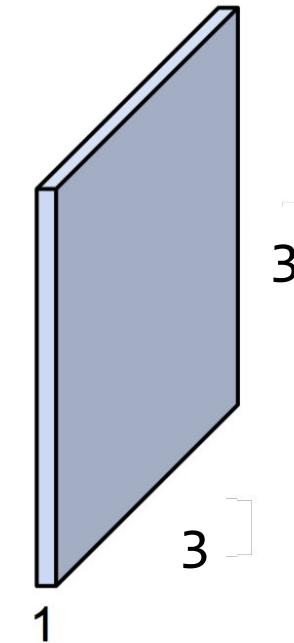
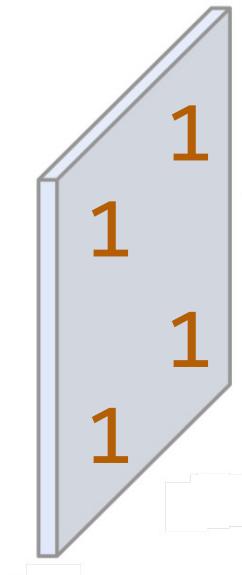
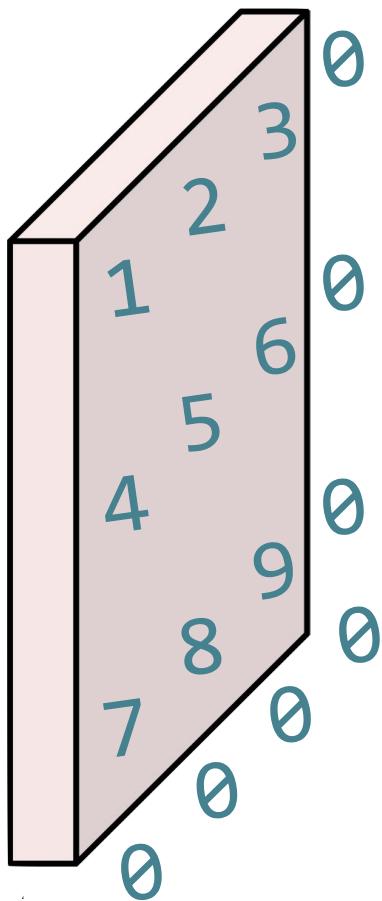
# Simple convolution layer

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, Padding: VALID

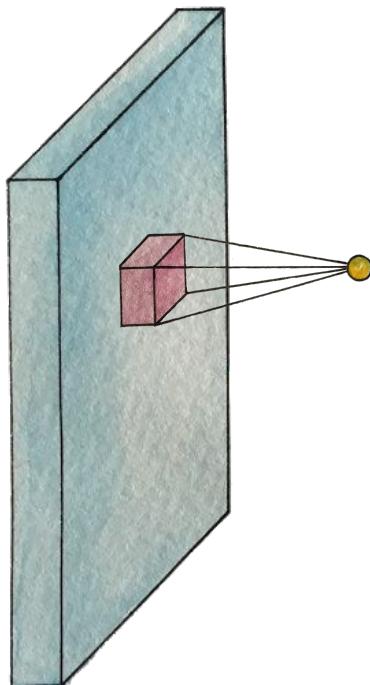


# Simple convolution layer

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, Padding: **SAME**



# Max pooling



Single depth slice

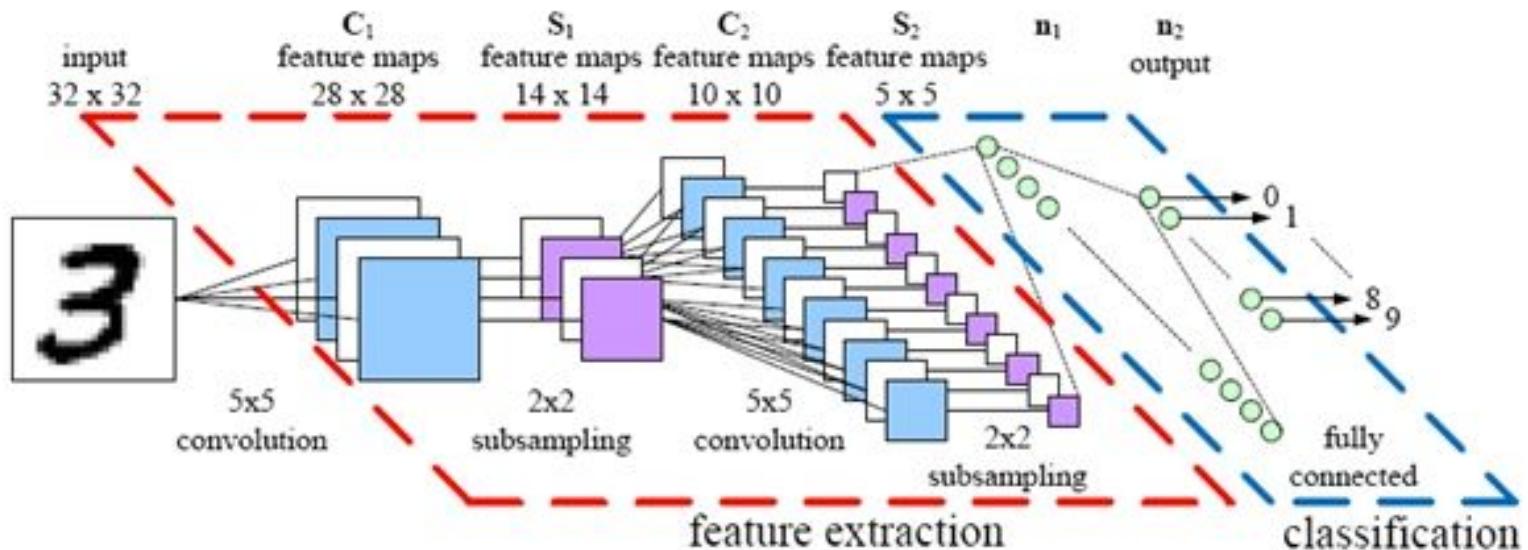
1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters  
and stride 2



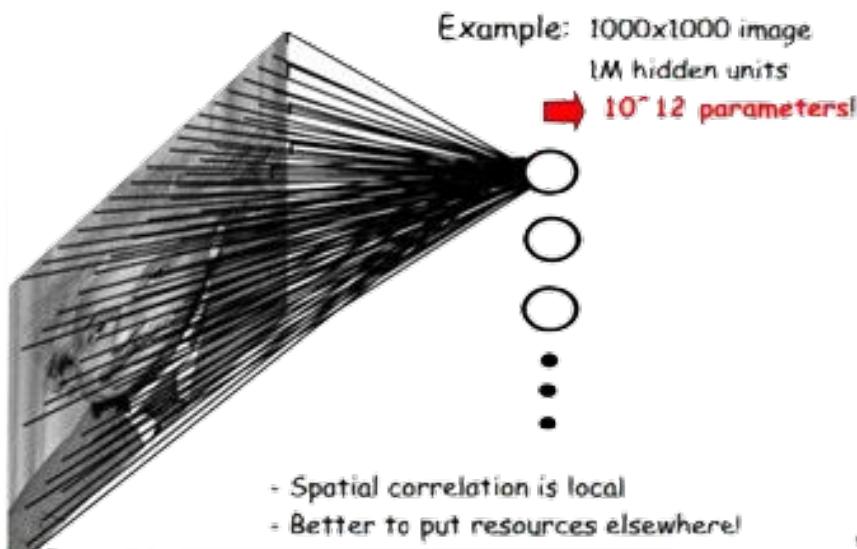
6	8
3	4

# CNN

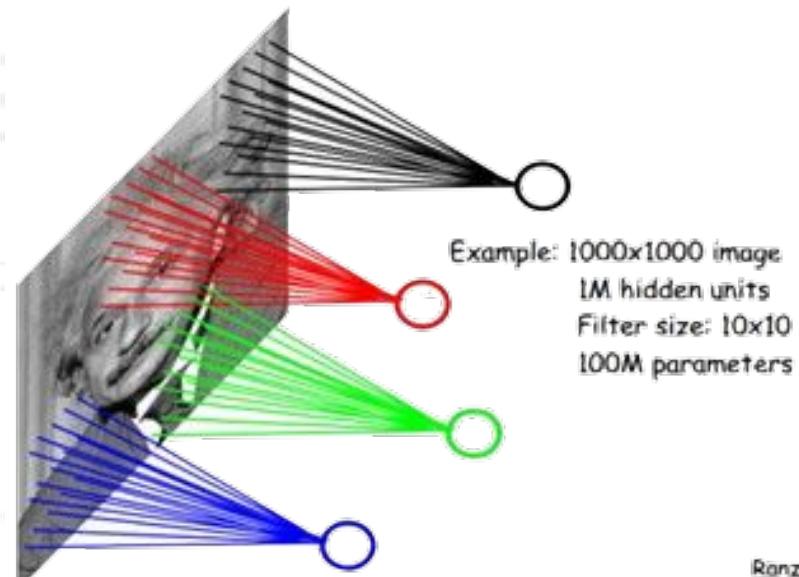


# Locally Connected Features

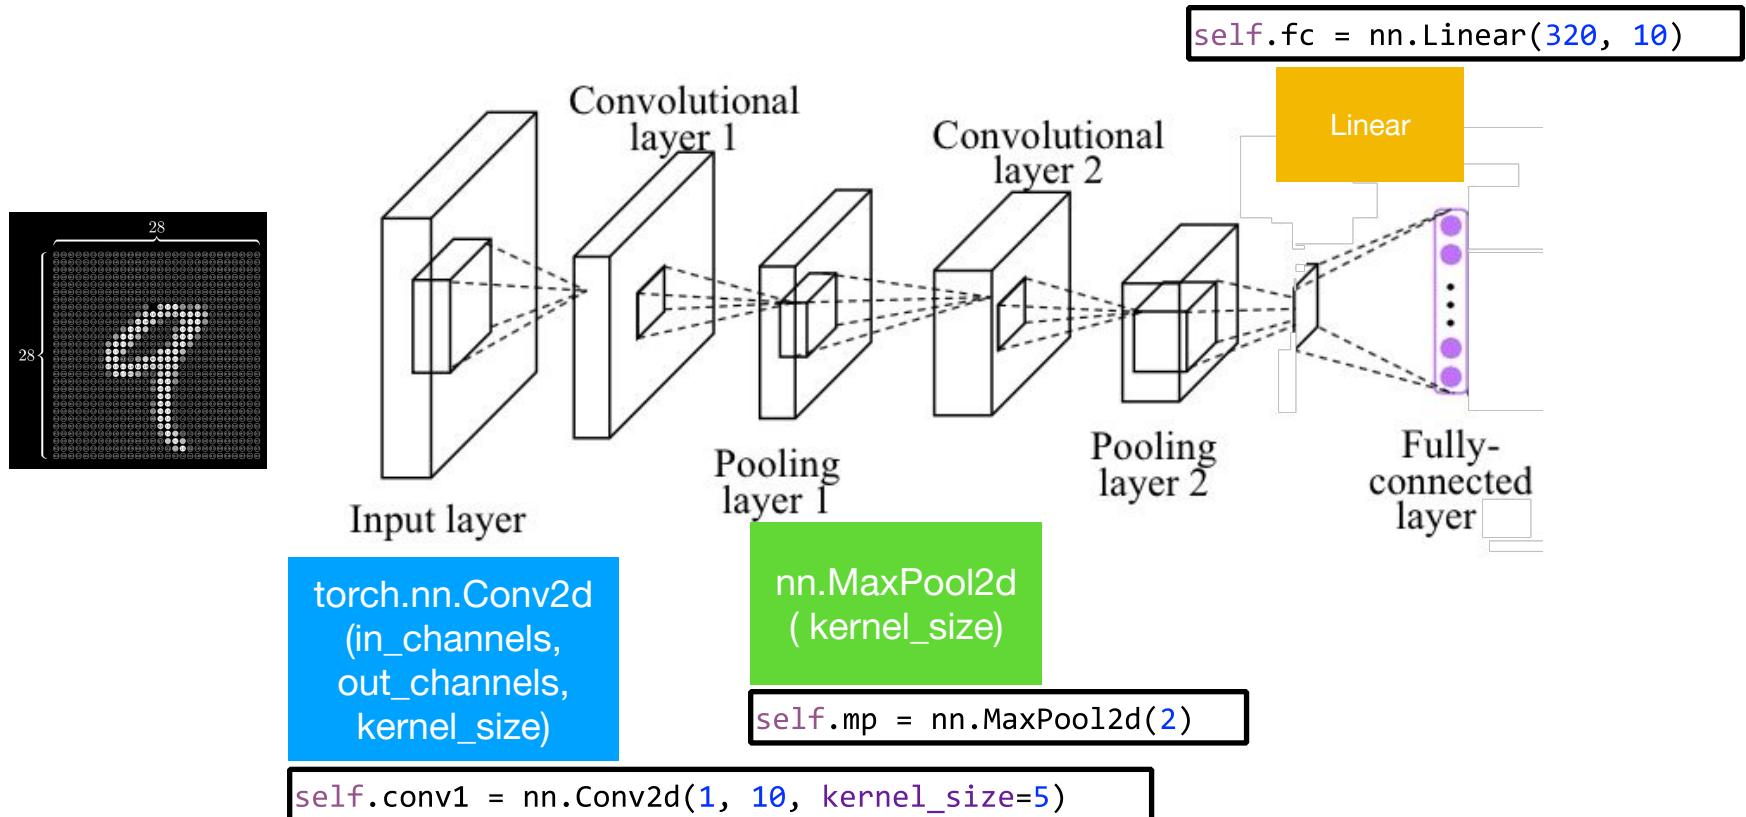
## FULLY CONNECTED NEURAL NET

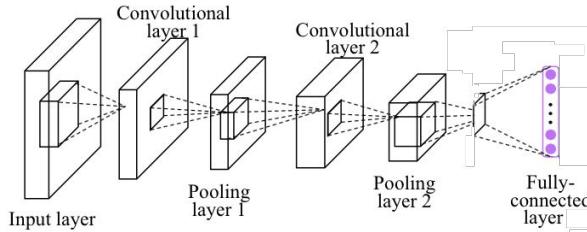


## LOCALLY CONNECTED NEURAL NET



# Simple CNN





# Simple CNN

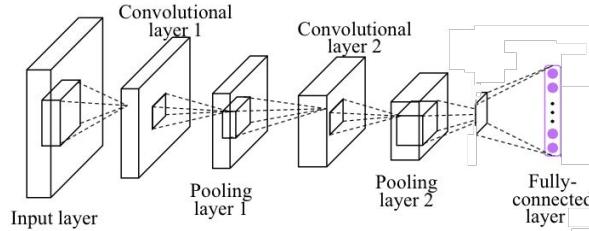


```

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(100???, 10) # ??? -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
  
```



# Simple CNN

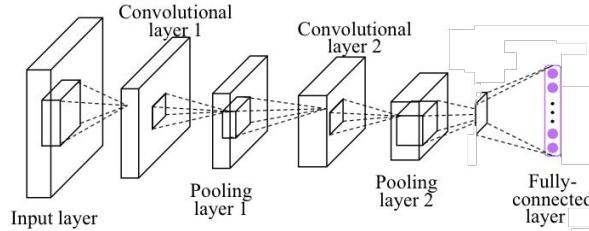


```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(100???, 10) # ??? -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

RuntimeError: size mismatch, m1: [64 x 320], m2: [100 x 10]



# Simple CNN

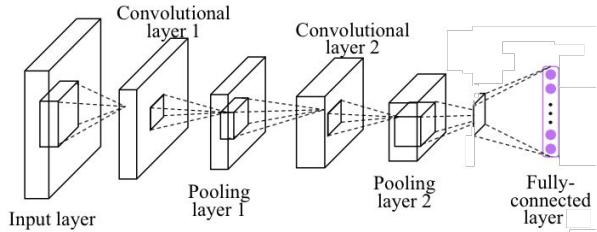


```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10) # 320 -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

RuntimeError: size mismatch, m1: [64 x 320], m2: [100 x 10]



# Simple CNN



```
class Net(nn.Module):

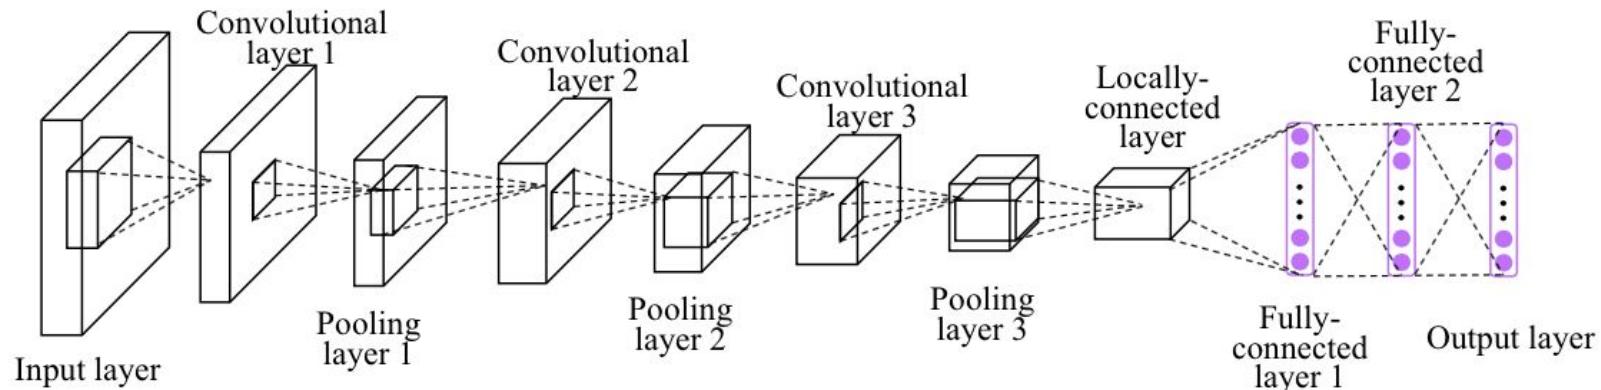
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10) # 320 -> 10

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
```

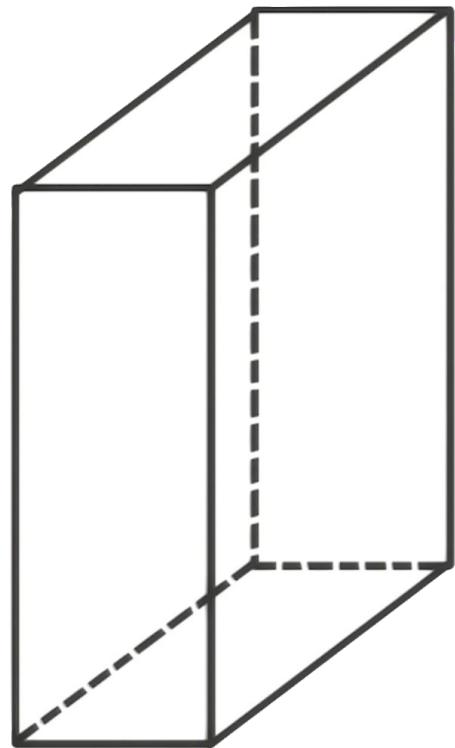
Train Epoch: 9 [46080/60000 (77%)]	Loss: 0.108415
Train Epoch: 9 [46720/60000 (78%)]	Loss: 0.140700
Train Epoch: 9 [47360/60000 (79%)]	Loss: 0.090830
Train Epoch: 9 [48000/60000 (80%)]	Loss: 0.031640
Train Epoch: 9 [48640/60000 (81%)]	Loss: 0.014934
Train Epoch: 9 [49280/60000 (82%)]	Loss: 0.090210
Train Epoch: 9 [49920/60000 (83%)]	Loss: 0.074975
Train Epoch: 9 [50560/60000 (84%)]	Loss: 0.058671
Train Epoch: 9 [51200/60000 (85%)]	Loss: 0.023464
Train Epoch: 9 [51840/60000 (86%)]	Loss: 0.018025
Train Epoch: 9 [52480/60000 (87%)]	Loss: 0.098865
Train Epoch: 9 [53120/60000 (88%)]	Loss: 0.013985
Train Epoch: 9 [53760/60000 (90%)]	Loss: 0.070476
Train Epoch: 9 [54400/60000 (91%)]	Loss: 0.065411
Train Epoch: 9 [55040/60000 (92%)]	Loss: 0.028783
Train Epoch: 9 [55680/60000 (93%)]	Loss: 0.008333
Train Epoch: 9 [56320/60000 (94%)]	Loss: 0.020412
Train Epoch: 9 [56960/60000 (95%)]	Loss: 0.036749
Train Epoch: 9 [57600/60000 (96%)]	Loss: 0.163087
Train Epoch: 9 [58240/60000 (97%)]	Loss: 0.117539
Train Epoch: 9 [58880/60000 (98%)]	Loss: 0.032256
Train Epoch: 9 [59520/60000 (99%)]	Loss: 0.026360

Test set: Average loss: 0.0483, Accuracy: 9846/10000 (98%)

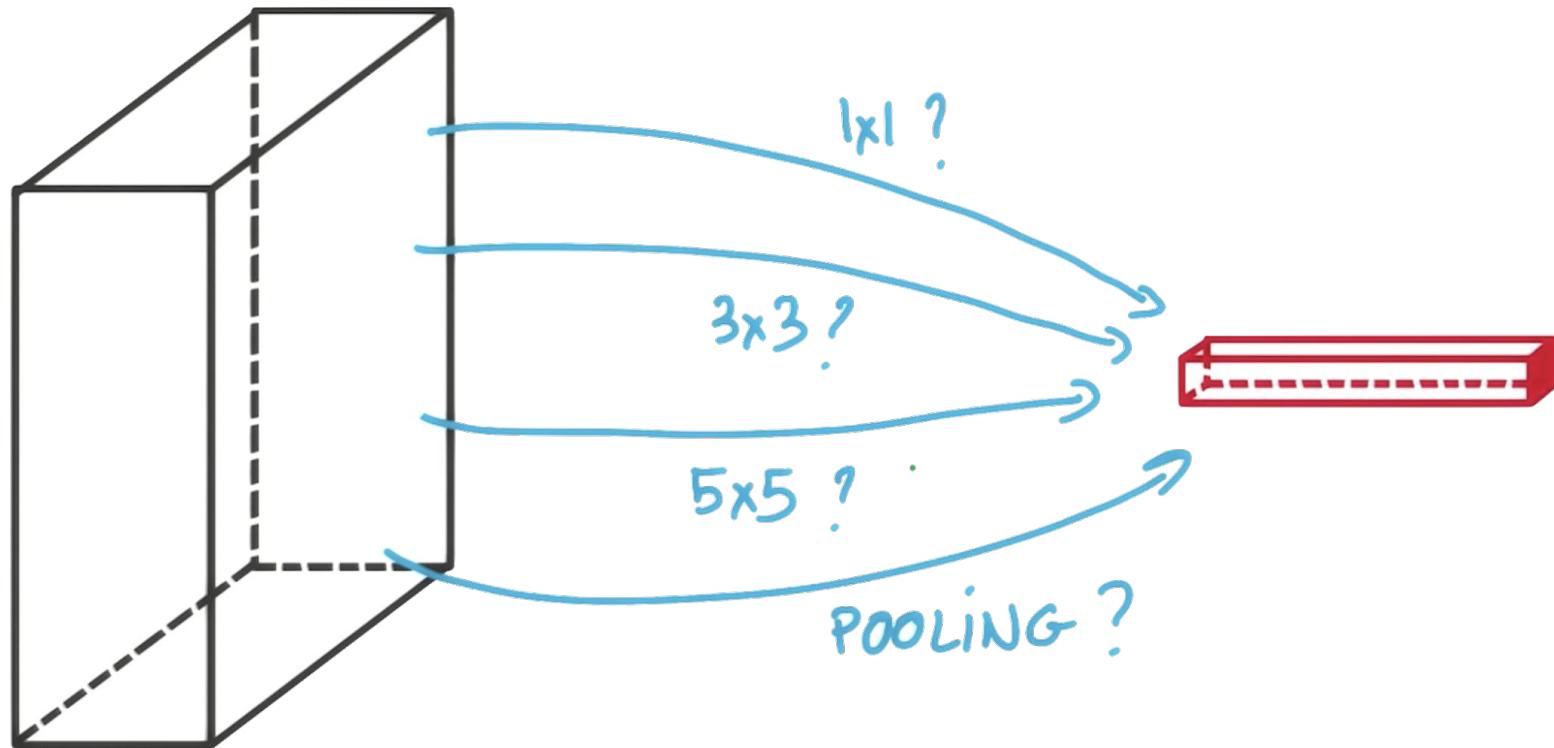
# Exercise 10-1: Implement CNN more layers



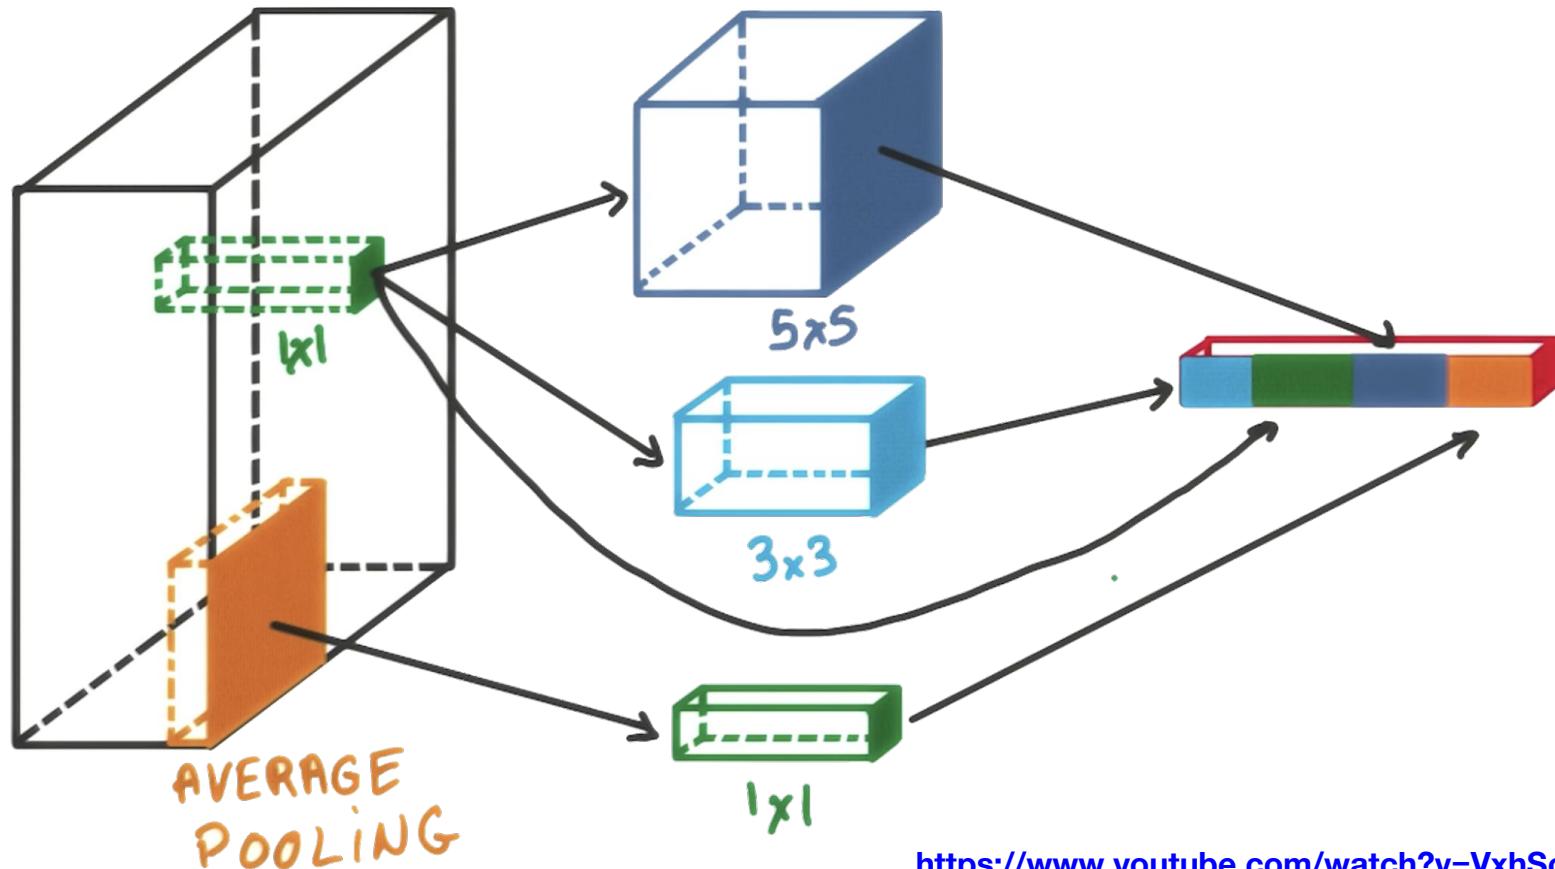
# INCEPTION MODULES



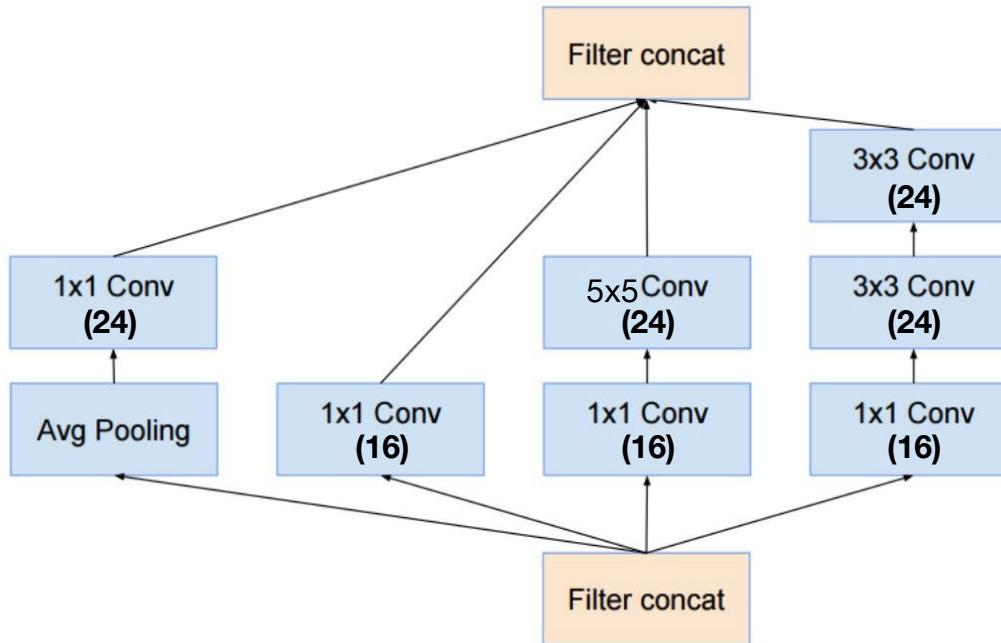
# INCEPTION MODULES

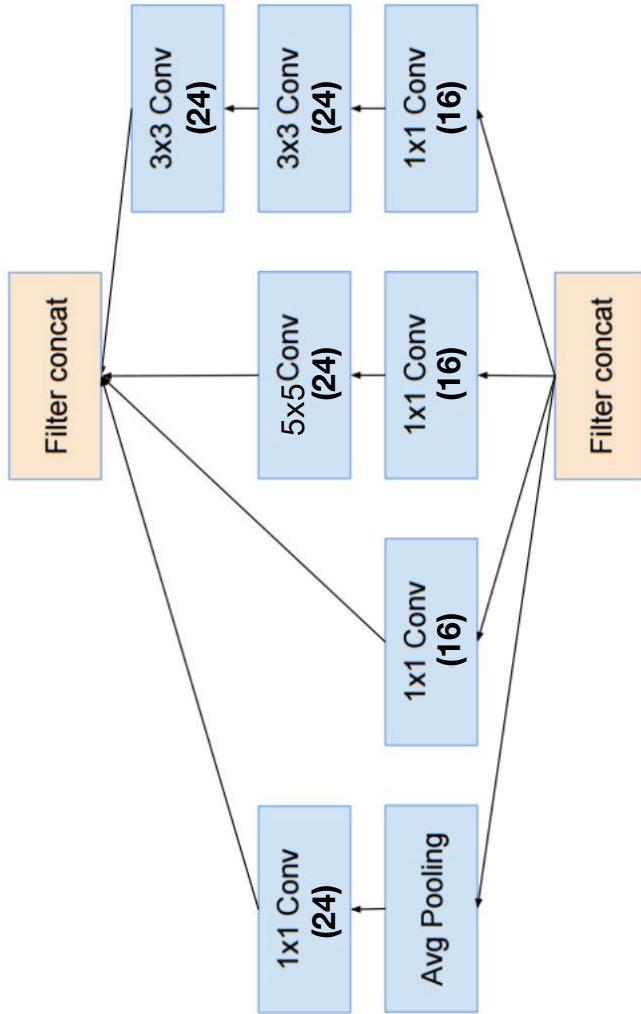


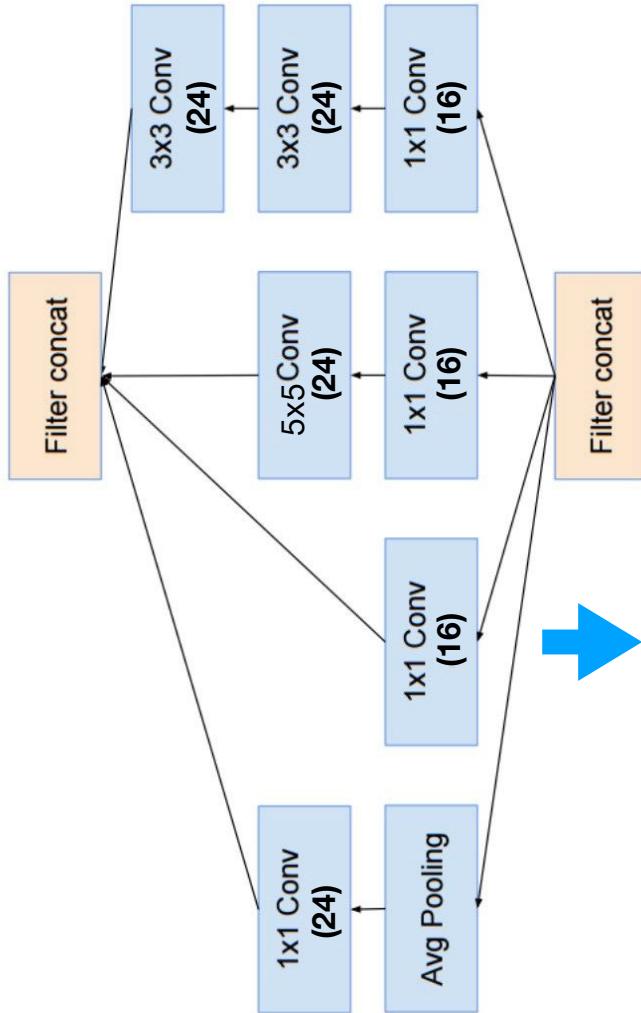
# INCEPTION MODULES



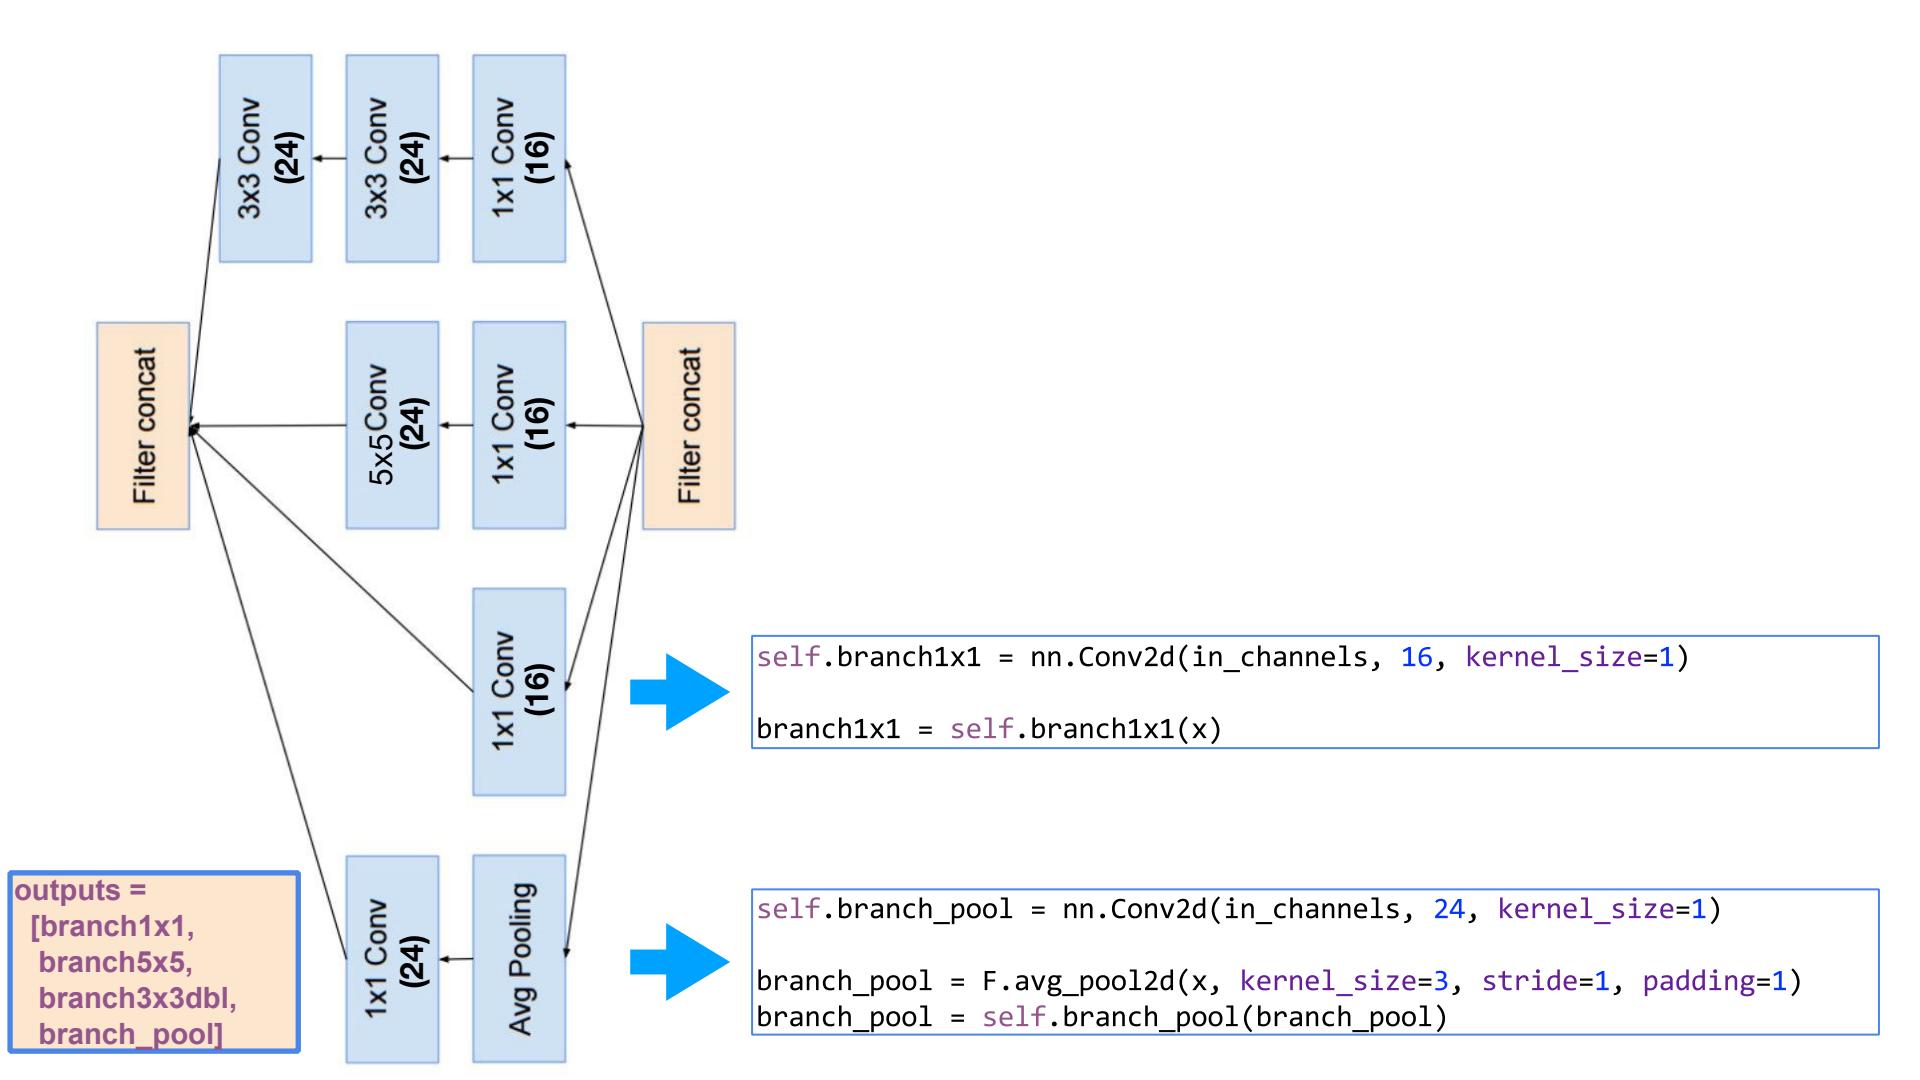
# Inception Module

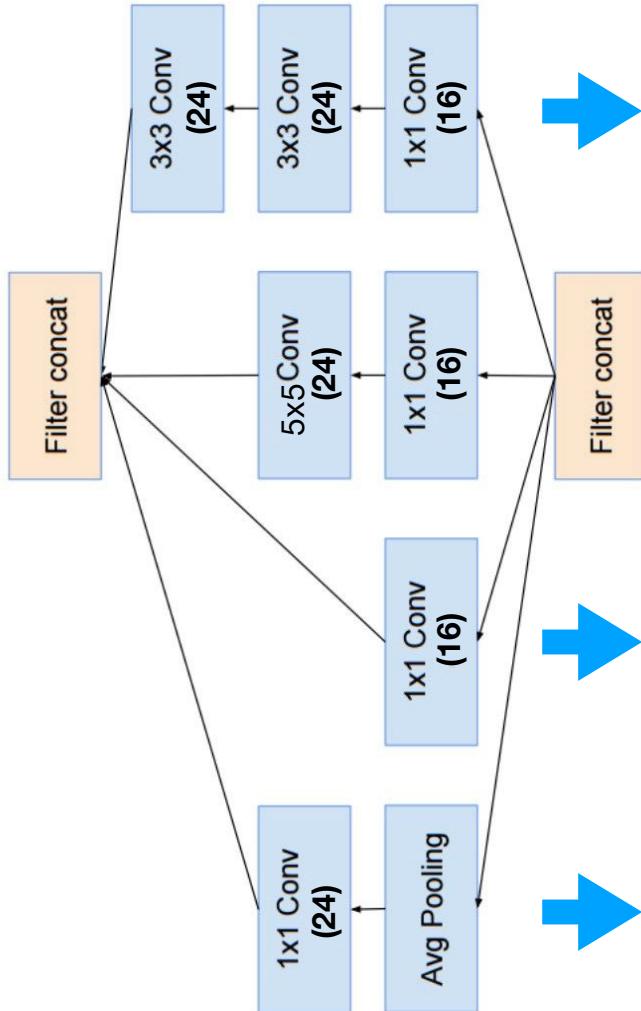






```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)  
branch1x1 = self.branch1x1(x)
```





```
self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)

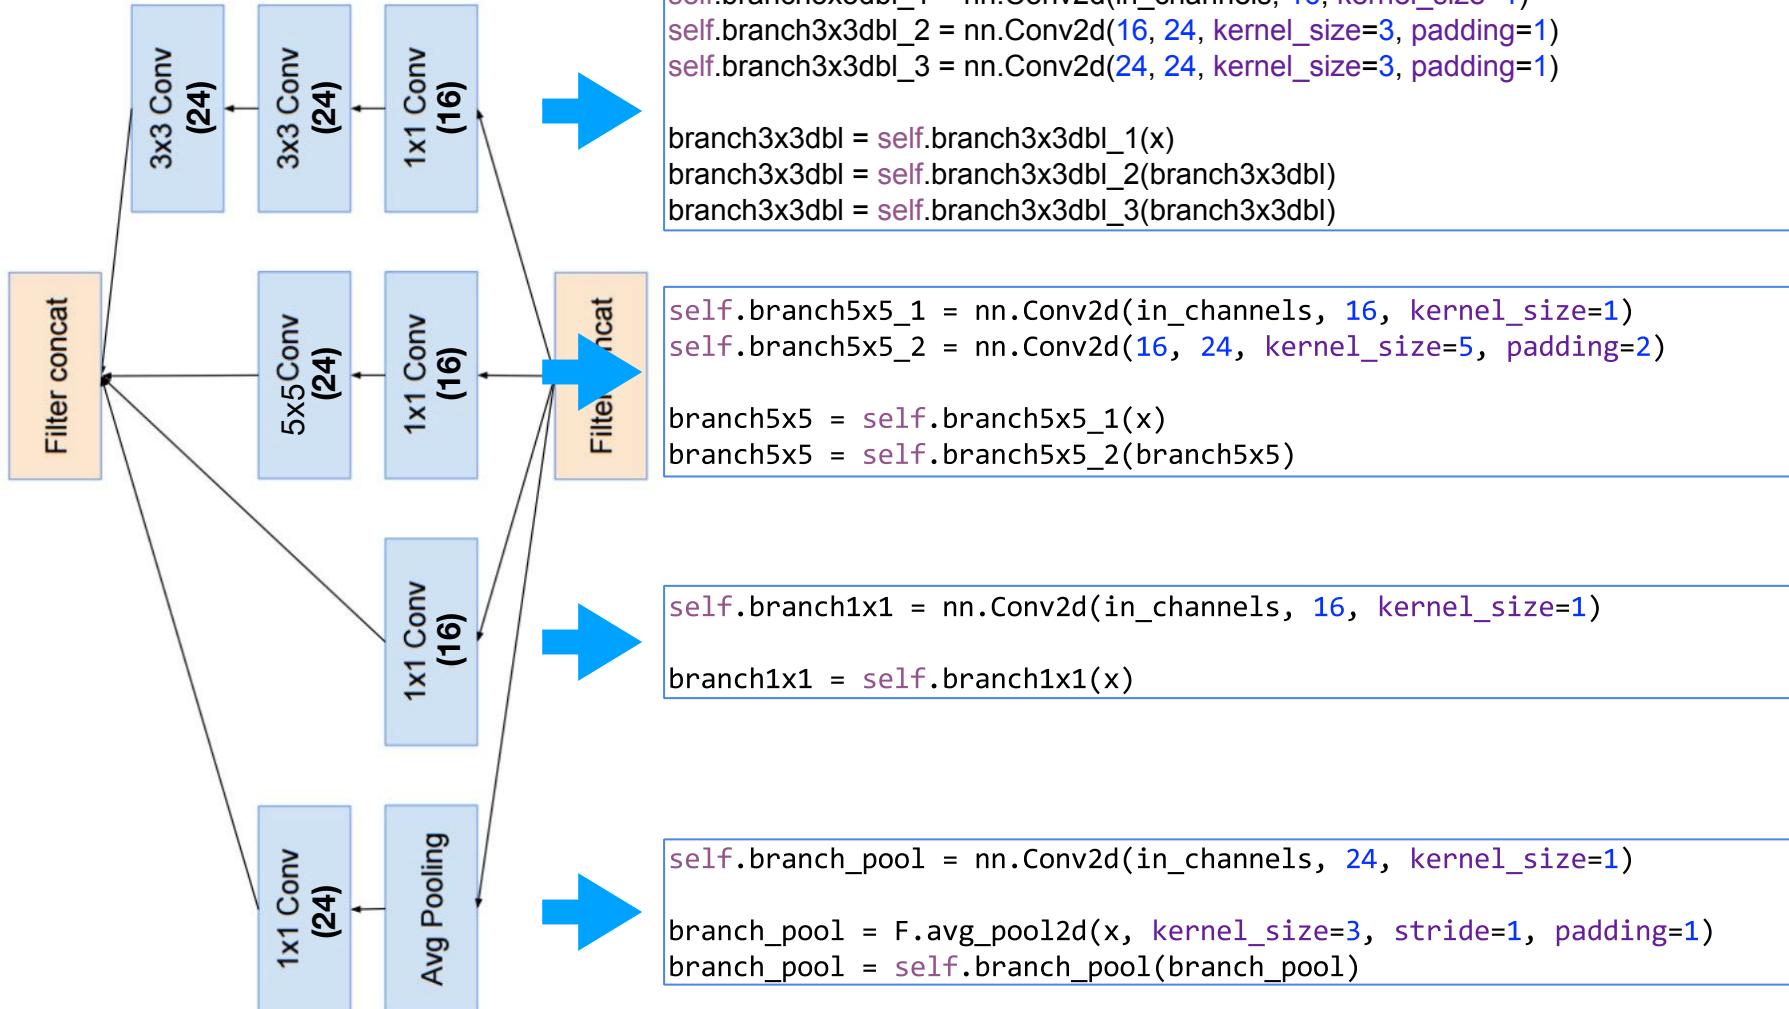
branch3x3dbl = self.branch3x3dbl_1(x)
branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)
```

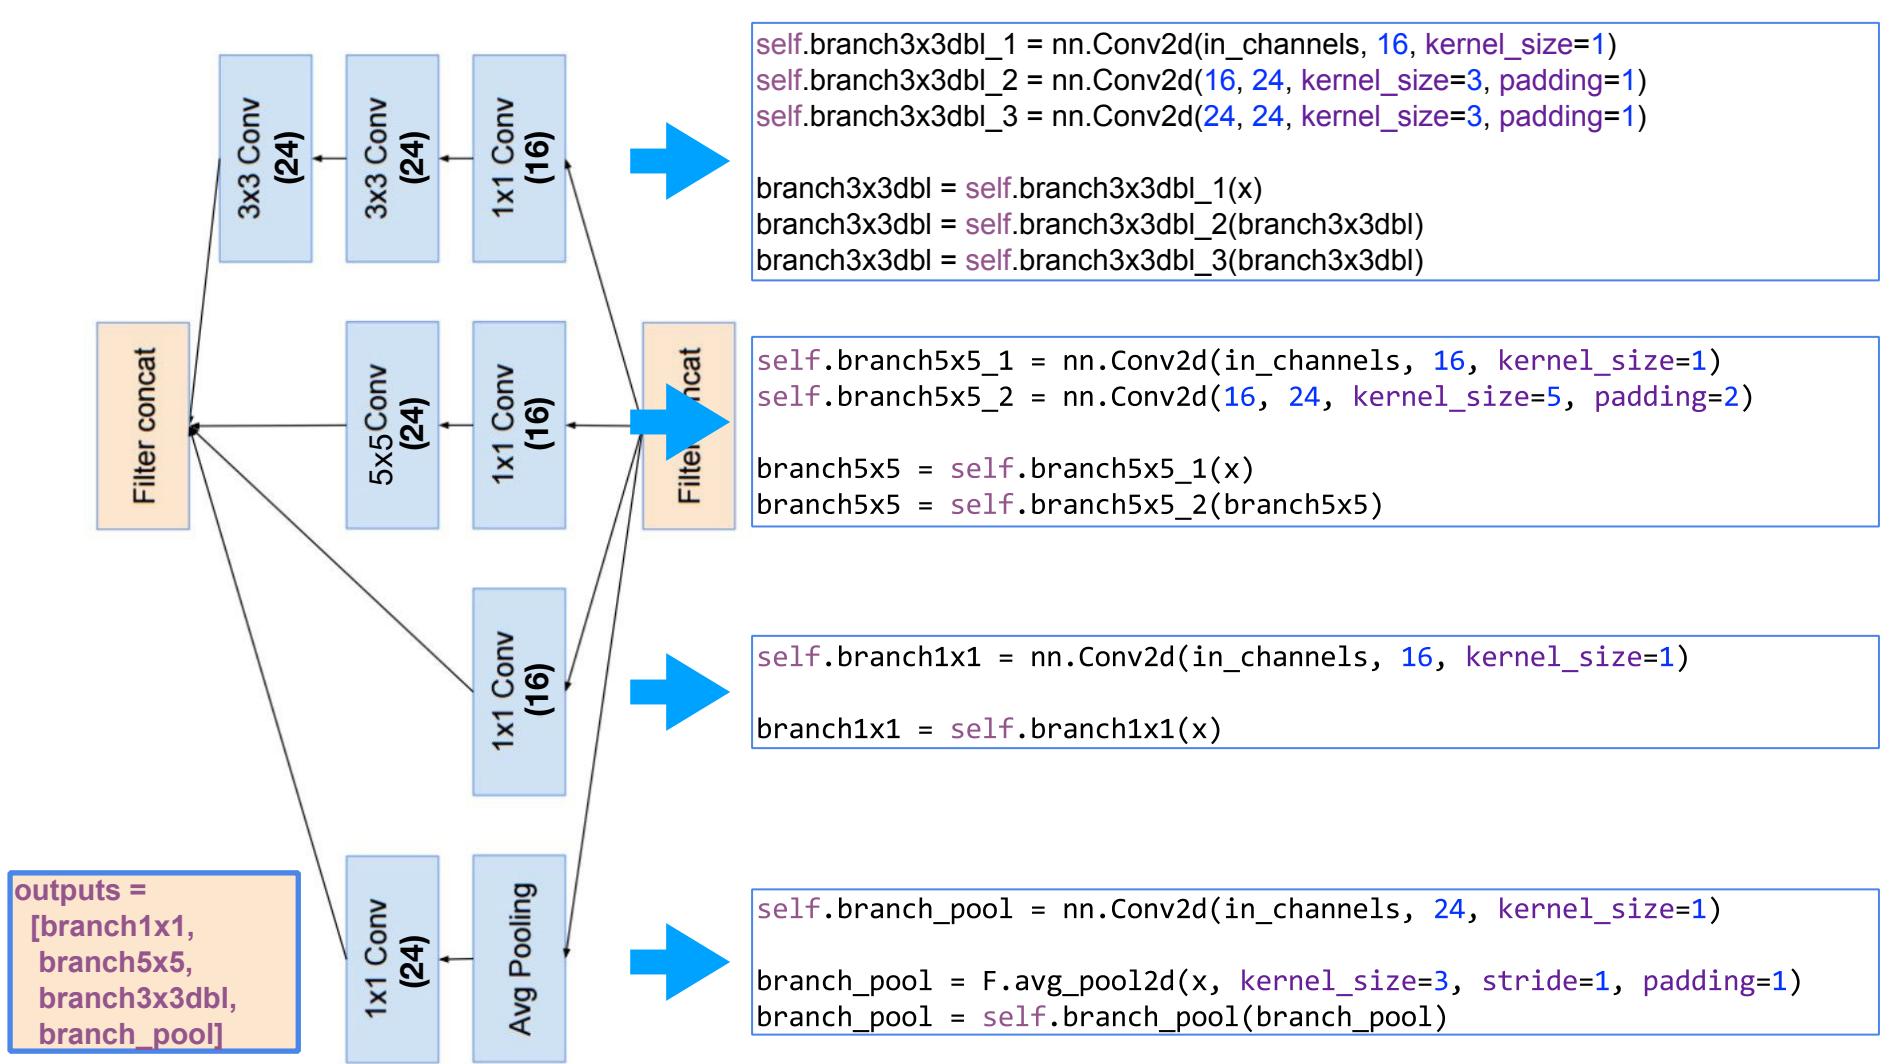
```
self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)

branch1x1 = self.branch1x1(x)
```

```
self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)

branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
branch_pool = self.branch_pool(branch_pool)
```





# Inception Module

```

class InceptionA(nn.Module):
    def __init__(self, in_channels):
        super(InceptionA, self).__init__()
        self.branch1x1 = nn.Conv2d(in_channels, 16, kernel_size=1)

        self.branch5x5_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
        self.branch5x5_2 = nn.Conv2d(16, 24, kernel_size=5, padding=2)

        self.branch3x3dbl_1 = nn.Conv2d(in_channels, 16, kernel_size=1)
        self.branch3x3dbl_2 = nn.Conv2d(16, 24, kernel_size=3, padding=1)
        self.branch3x3dbl_3 = nn.Conv2d(24, 24, kernel_size=3, padding=1)

        self.branch_pool = nn.Conv2d(in_channels, 24, kernel_size=1)

    def forward(self, x):
        branch1x1 = self.branch1x1(x)

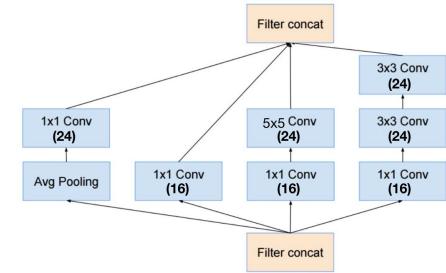
        branch5x5 = self.branch5x5_1(x)
        branch5x5 = self.branch5x5_2(branch5x5)

        branch3x3dbl = self.branch3x3dbl_1(x)
        branch3x3dbl = self.branch3x3dbl_2(branch3x3dbl)
        branch3x3dbl = self.branch3x3dbl_3(branch3x3dbl)

        branch_pool = F.avg_pool2d(x, kernel_size=3, stride=1, padding=1)
        branch_pool = self.branch_pool(branch_pool)

        outputs = [branch1x1, branch5x5, branch3x3dbl, branch_pool]
        return torch.cat(outputs, 1)

```



```

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(88, 20, kernel_size=5)

        self.incept1 = InceptionA(in_channels=10)
        self.incept2 = InceptionA(in_channels=20)

        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(1408, 10)

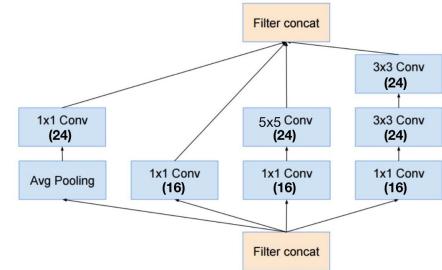
    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = self.incept1(x)
        x = F.relu(self.mp(self.conv2(x)))
        x = self.incept2(x)
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)

```

# Inception Module

Train Epoch: 9 [44800/60000 (75%)]	Loss: 0.064180
Train Epoch: 9 [45440/60000 (76%)]	Loss: 0.020339
Train Epoch: 9 [46080/60000 (77%)]	Loss: 0.061476
Train Epoch: 9 [46720/60000 (78%)]	Loss: 0.039662
Train Epoch: 9 [47360/60000 (79%)]	Loss: 0.026798
Train Epoch: 9 [48000/60000 (80%)]	Loss: 0.071569
Train Epoch: 9 [48640/60000 (81%)]	Loss: 0.003835
Train Epoch: 9 [49280/60000 (82%)]	Loss: 0.005564
Train Epoch: 9 [49920/60000 (83%)]	Loss: 0.020116
Train Epoch: 9 [50560/60000 (84%)]	Loss: 0.128114
Train Epoch: 9 [51200/60000 (85%)]	Loss: 0.016599
Train Epoch: 9 [51840/60000 (86%)]	Loss: 0.006995
Train Epoch: 9 [52480/60000 (87%)]	Loss: 0.111267
Train Epoch: 9 [53120/60000 (88%)]	Loss: 0.052126
Train Epoch: 9 [53760/60000 (90%)]	Loss: 0.034962
Train Epoch: 9 [54400/60000 (91%)]	Loss: 0.029465
Train Epoch: 9 [55040/60000 (92%)]	Loss: 0.031482
Train Epoch: 9 [55680/60000 (93%)]	Loss: 0.015132
Train Epoch: 9 [56320/60000 (94%)]	Loss: 0.010435
Train Epoch: 9 [56960/60000 (95%)]	Loss: 0.014344
Train Epoch: 9 [57600/60000 (96%)]	Loss: 0.014952
Train Epoch: 9 [58240/60000 (97%)]	Loss: 0.153132
Train Epoch: 9 [58880/60000 (98%)]	Loss: 0.112024
Train Epoch: 9 [59520/60000 (99%)]	Loss: 0.009406

Test set: Average loss: 0.0470, Accuracy: 9866/10000 (99%)



```

class Net(nn.Module):

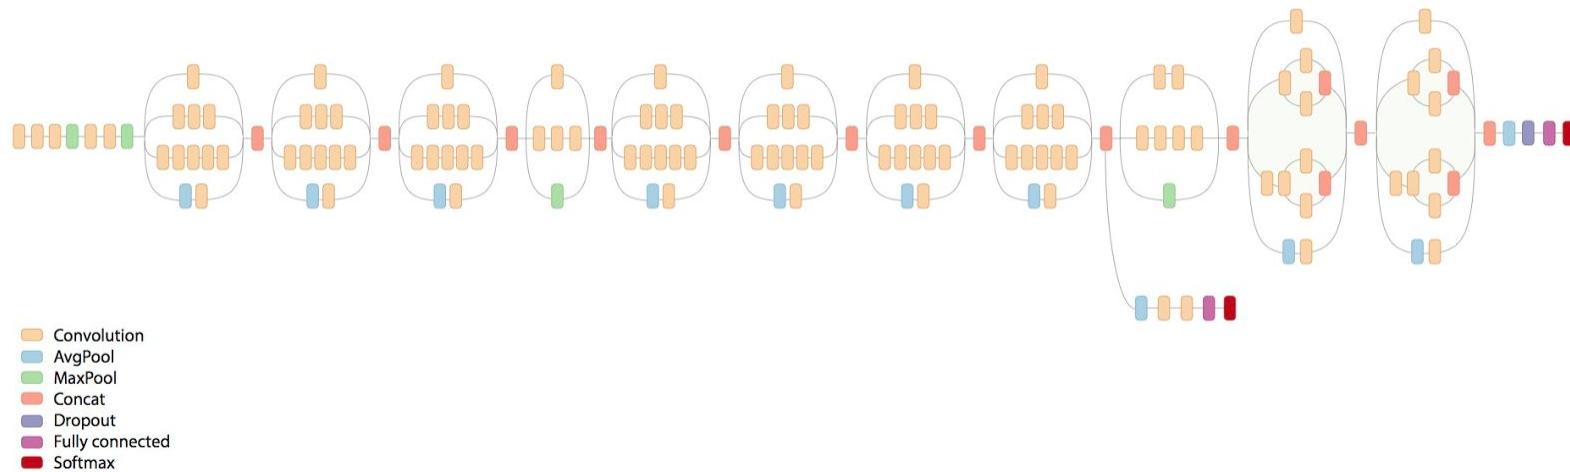
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(88, 20, kernel_size=5)

        self.incept1 = InceptionA(in_channels=10)
        self.incept2 = InceptionA(in_channels=20)

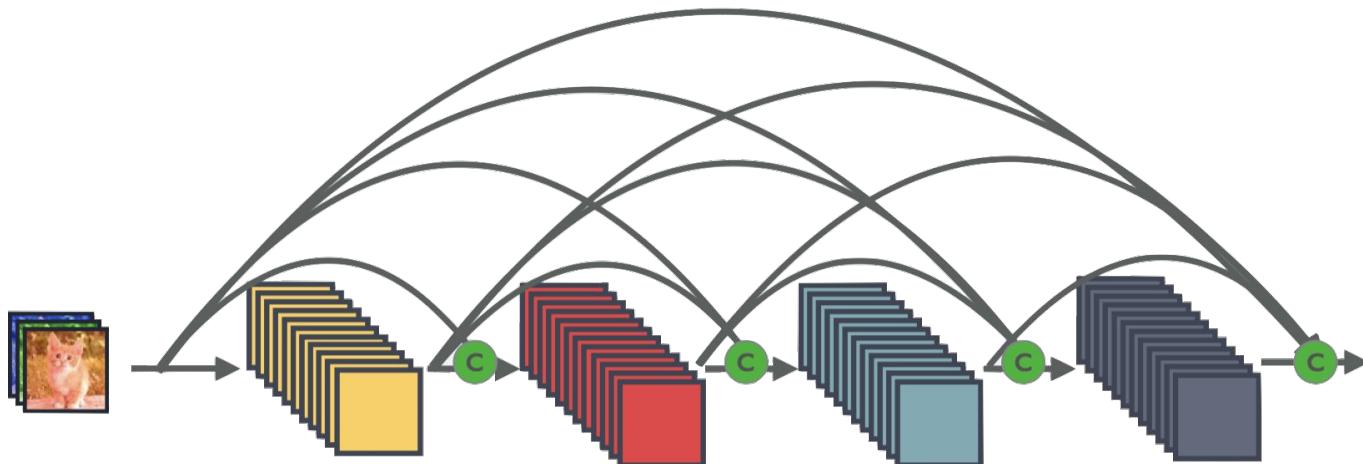
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(1408, 10)

    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = self.incept1(x)
        x = F.relu(self.mp(self.conv2(x)))
        x = self.incept2(x)
        x = x.view(in_size, -1) # flatten the tensor
        x = self.fc(x)
        return F.log_softmax(x)
  
```

# Exercise 10-2: Implement full inception v3/v4



# Exercise 10-3: Implement DenseNet





**WHAT  
NEXT?**

A woman with long brown hair tied back in a ponytail, wearing a dark blue blazer over a light blue shirt, is shown in profile facing right. She has her right hand raised to her forehead, with her fingers partially covering her eyes as if peering into the distance or shielding her eyes from bright light. Her gaze is directed upwards and to the right. In the upper right corner of the image, there is a graphic element consisting of the words "WHAT NEXT?" in large, bold, sans-serif font. The word "WHAT" is in orange and "NEXT?" is in blue. A simple line drawing of a lit lightbulb is positioned to the right of the question mark.

## Lecture 11: RNN